

Author: Guilherme Indicciate - AI Engineer - Full-Stack for Generative AI Applications  
Delivered on: January 8th, 2026



# Technical Challenge – AI Engineer Junior

## Artefact Assistant | Data & AI to Drive Impact

An intelligent AI assistant that automatically decides when to use external tools (calculator, FX, crypto) or respond using its own knowledge.

The assistant supports **tool routing**, **multi-tool chaining**, **optional conversational memory**, and an **interactive CLI**, implemented using **LangGraph** and **LangChain**.

## Project Structure

### Backend

```
backend/
└── src/
    ├── llm/
    │   ├── models/
    │   │   └── models.py      # Typed agent state definition
    │   ├── assistant.py      # Assistant logic (with and without memory)
    │   ├── graph.py          # LangGraph execution graph
    │   └── llm_call.py       # LLM invocation and configuration
    ├── observability/
    │   └── context.py        # Tracing and tool-usage context
    └── tools/
        └── tools.py          # External tools (calculator, FX, crypto)
    ├── .env
    ├── main.py              # Environment variables (API keys, config)
    └── requirements.txt      # Backend entry point (FastAPI / CLI)
                            # Python dependencies
```

### Frontend

```
frontend/
└── src/
    └── app/
        └── page.tsx          # Main application entry point (state + layout
    orchestration)
            ├── page.module.css # Global page layout and container styles
            └── chat.module.css  # Chat area, message list, scrolling, empty state
    styles
```

```
|   └── controls.module.css # Input field, buttons, and interaction controls  
styles  
    └── globals.css          # Global CSS reset, fonts, and base styles  
  
components/  
  ├── Header.tsx           # Header with branding (logo + title)  
  ├── ChatBox.tsx          # Scrollable chat container with auto-scroll logic  
  ├── MessageBubble.tsx    # Single message renderer (user vs assistant)  
  └── TypingIndicator.tsx  # Animated typing indicator while LLM responds  
  
lib/  
  ├── api.ts                # Backend API client (chat, reset, etc.)  
  └── session.ts            # Session ID management for conversational memory
```

## 🚀 How to Run

### 1. Installation

```
git clone https://github.com/guiindicate/Artefact_Technical_Challenge.git  
cd source/backend/src  
pip install -r requirements.txt
```

### 2. Configuration

```
# Create a ` `.env` file with your OpenAI API key:  
OPENAI_API_KEY=sk-proj-xxxxxxxxxxxxxx
```

#### How to obtain an OpenAI API key:

- Go to <https://platform.openai.com/api-keys>
- Create a new API key
- Copy it into the `.env` file, using `OPENAI_API_KEY` variable

### 3. Run the Backend

```
uvicorn main:app --port 8001 --reload
```

### 4. Run the Frontend

```
# Navigate to the Frontend Directory  
cd source/frontend
```

```
# Create a `*.env.local` file and configure the backend API base URL  
NEXT_PUBLIC_API_BASE=http://localhost:8001
```


```
# Install next.js dependencies  
npm install
```

```
# Activate frontend application  
npm run dev
```

Now, you can open the application in <https://localhost:3000/>

## 🧠 Architecture

### LangGraph Execution Flow



## Core Components

### ./tools/

Defines all **external tools** available to the assistant. Each tool is deterministic, isolated, and observable.

- `calculator()` Performs deterministic arithmetic operations (used for exact calculations and post-processing).
- `fx_convert()` Converts fiat currencies using a public FX API, returning precise exchange rates.
- `crypto_convert()` Retrieves cryptocurrency prices and conversions using the CoinGecko public API (no API key required).

Each tool explicitly reports when it is used, enabling clear observability of routing decisions.

---

### ./llm/llm\_call.py

LLM invocation layer with tool binding.

- `call_model(state: AgentState) -> dict` Invokes the language model (`ChatOpenAI`) with tool calling enabled. The function binds all available tools (`calculator`, `fx_convert`, `crypto_convert`) to the model, sends the current message history, and returns the LLM response wrapped in the expected graph state format.

This component keeps model invocation logic isolated from graph orchestration, making it easy to swap models, adjust parameters, or extend tool bindings without impacting the execution flow.

---

### ./graph/

Defines the **LangGraph execution graph** and the agent control flow.

- `should_continue()` Conditional routing logic that checks for `tool_calls` and decides whether to:
    - route execution to the Tools node, or
    - terminate the graph.
  - `create_graph()` Assembles and compiles the LangGraph workflow, connecting agent and tool nodes into a deterministic loop.
- 

### ./assistant/

High-level assistant interface layer.

- `AssistantWithMemory` A thin wrapper around the compiled LangGraph workflow that maintains an in-memory `conversation_history` per instance. On each turn, it appends the new `HumanMessage`, runs the graph with the full history, stores the updated message list, and returns the latest `AIMessage`.
- `chat(query: str) -> str` Main entry point for multi-turn interactions. Sends the full conversation context into the graph and returns the assistant's latest reply.

- `clear_history() -> None` Resets the conversation context for a fresh session.
- 

## ./observability/ (or equivalent telemetry utilities)

Provides **lightweight observability** for runtime execution.

- Tracks **which tools were used** during a request
- Attaches a **trace\_id** to each interaction
- Preserves **context propagation** across turns when memory is enabled

This makes agent behavior transparent, debuggable, and aligned with production-grade AI system expectations.

---

## 💡 How It Works

1. **User submits a question** -> converted to a `HumanMessage`
  2. **Agent node (LLM)** analyzes the context
  3. **Routing decision:**
    - Math-related -> calls `calculator`
    - Currency conversion -> calls `fx_convert`
    - Crypto price/conversion -> calls `crypto_convert`
    - General knowledge -> answers directly
  4. **Tools node** executes the requested tool
  5. **Agent loops back** to generate the final response
  6. **Conversation memory** (optional) preserves context across turns
- 

## 🌐 Implementation Logic

- Built a **minimal, deterministic agent loop** using LangGraph
- Used an **Agent (LLM) node** and a **Tools node**
- Enabled **tool calling** via `bind_tools`, allowing the model to decide when to invoke external functions
- Routed execution based on the presence of `tool_calls` in the LLM response
- Continued looping until no further tool execution was required
- Used **deterministic tools** for calculations and conversions to ensure accuracy and trust
- Added **observability logs** to clearly show when tools are used
- Implemented **memory as a thin wrapper**, without modifying the core execution graph

## User Experience & Product Considerations

Beyond the core technical requirements, several UX-oriented decisions were intentionally added to make the assistant feel closer to a real product:

- Typing indicator while the LLM is processing (Frontend), reducing perceived latency
- Clear visibility in logs of when external tools are used
- Deterministic tool outputs for reliability and explainability
- Optional conversational memory for natural multi-turn interactions

These small details improve usability while keeping the system simple and maintainable.

## Conversational Memory (Optional Enhancement)

Although **conversational memory was not a required feature** for this challenge, it was included as an optional enhancement to demonstrate how the assistant could evolve in a real-world scenario.

To keep the core solution aligned with the challenge scope:

- The **LangGraph execution graph remains stateless**
- Memory is implemented as a **thin wrapper ([AssistantWithMemory](#))**
- The same graph can be reused in **single-shot** or **conversational** modes

This preserves clarity and extensibility without increasing system complexity.

## Examples

```
>>> chat("What does Artefact do in a few words?")
"Artefact is a data and digital consulting company that helps organizations
accelerate AI and data adoption to drive business impact."

>>> chat("How much is 0.1 BTC in BRL?")
"[TOOL] Using CRYPTO Converter (CoinGecko)"
"🤖 Artefact Assistant: 0.1 BTC is approximately 49,031.40 BRL."

>>> chat("How much is 1 USD in BRL?")
"[TOOL] Using FX CONVERTER"
"🤖 Artefact Assistant: 1 USD is approximately 5.39 BRL.

>>> chat("Using the current BTC price in BRL and the USD/BRL exchange rate, what
is the price of 0.1 BTC in USD?")
"[TOOL] Using CRYPTO CONVERTER (CoinGecko)"
"[TOOL] Using FX CONVERTER"
"[TOOL] Using local CALCULATOR"
"🤖 Artefact Assistant: The price of 0.1 BTC is approximately 9,103.65 USD."
```

These examples demonstrate how the assistant **combines multiple tools across turns** while preserving conversational context. The assistant retrieves external data, performs deterministic calculations, and chains results logically to solve multi-step user requests in a transparent and explainable way.

## ⌚ Learnings & Next Steps

### What I learned

This project strengthened my understanding of how to orchestrate LLMs with external tools using execution graphs, and how thoughtful decisions around state management and routing directly impact an assistant's clarity, reliability, and scalability.

### What I'd do with more time

With additional time, I would focus on making routing more robust, expanding test coverage and

observability, and improving how conversational memory is managed over longer sessions. In particular, I would introduce memory pruning strategies (such as limiting context length or summarizing earlier interactions) to control token usage and costs while preserving relevant context.

---

## Technologies

- **LangGraph** – agent orchestration and control flow
  - **LangChain** – tools and LLM integration
  - **OpenAI GPT-4o-mini** – language model
  - **FastAPI** – backend API for chat sessions and tool execution
  - **Next.js (React)** – interactive frontend chat interface
  - **Python 3.9+**
- 

## References

- <https://langchain-ai.github.io/langgraph/>
  - <https://python.langchain.com/docs/modules/tools/>
  - <https://platform.openai.com/docs/guides/function-calling>
-